

LABWORK 3

HELLO CUDA

TRAN Quy Ban
Department of ICT

Introduction

Implement code for labwork 1 using CUDA.

How to implement the labwork

```
__global__ void grayscale(uchar3 *input, uchar3 *output) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    output[tid].x = (input[tid].x + input[tid].y + input[tid].z) / 3;
    output[tid].z = output[tid].y = output[tid].x;
}
```

```
void Labwork::labwork3_GPU() {
    uchar3 *devInput;
    uchar3 *devGray;
    uchar3 *hostGray;
    int pixelCount = inputImage->width * inputImage->height;
    char *hostInput = static_cast<char *>(malloc(pixelCount * 3));
    // Allocate CUDA memory
    cudaMalloc(&devInput, pixelCount * sizeof(uchar3));
    cudaMalloc(&devGray, pixelCount * sizeof(float));
    // Copy CUDA Memory from CPU to GPU
    cudaMemcpy(devInput, hostInput, pixelCount * sizeof(uchar3), cudaMemcpyHostToDevice);
    // Processing
    int blockSize = 64;
    int numBlock = pixelCount / blockSize;
    grayscale<<<numBlock, blockSize>>>>(devInput, devGray);

    // Copy CUDA Memory from GPU to CPU
    cudaMemcpy(hostGray, devGray, pixelCount * sizeof(float), cudaMemcpyDeviceToHost);
    // Cleaning
    cudaFree(devInput);
}
```

The result of the labwork is shown as below:

What is the speed up?

Try experimenting with different block size values

Plot a graph of block size vs speedup

Discuss